



WIZELINE ACADEMY

Grow your career:
Free courses in Artificial Intelligence,
Software Development, User Experience and
More

WIFI: WizelineAcademy
Password: academyGDL
Slack Channel:



@WizelineAcademy



/WizelineAcademy



academy.wizeline.com



Get notified about courses:
tinyurl.com/WL-academy



Big Data Engineering with Spark

RDDs, Dataframes
and Datasets

A red notebook with a grid pattern is open on a wooden desk. A black pen and a gold smartphone are resting on the notebook. The text "In this episode..." is written in white on the notebook's cover.

In this episode...

- **Unstructured versus Structured Data**

- **Dataframes**

- **Datasets**



By the end of this session, you'll be able to explain:

- **Differences between unstructured and structured data**
- **How to use Dataframes**
- **How to use Datasets**
- **When to use Resilient Distributed Datasets (RDDs), Dataframes or Datasets**

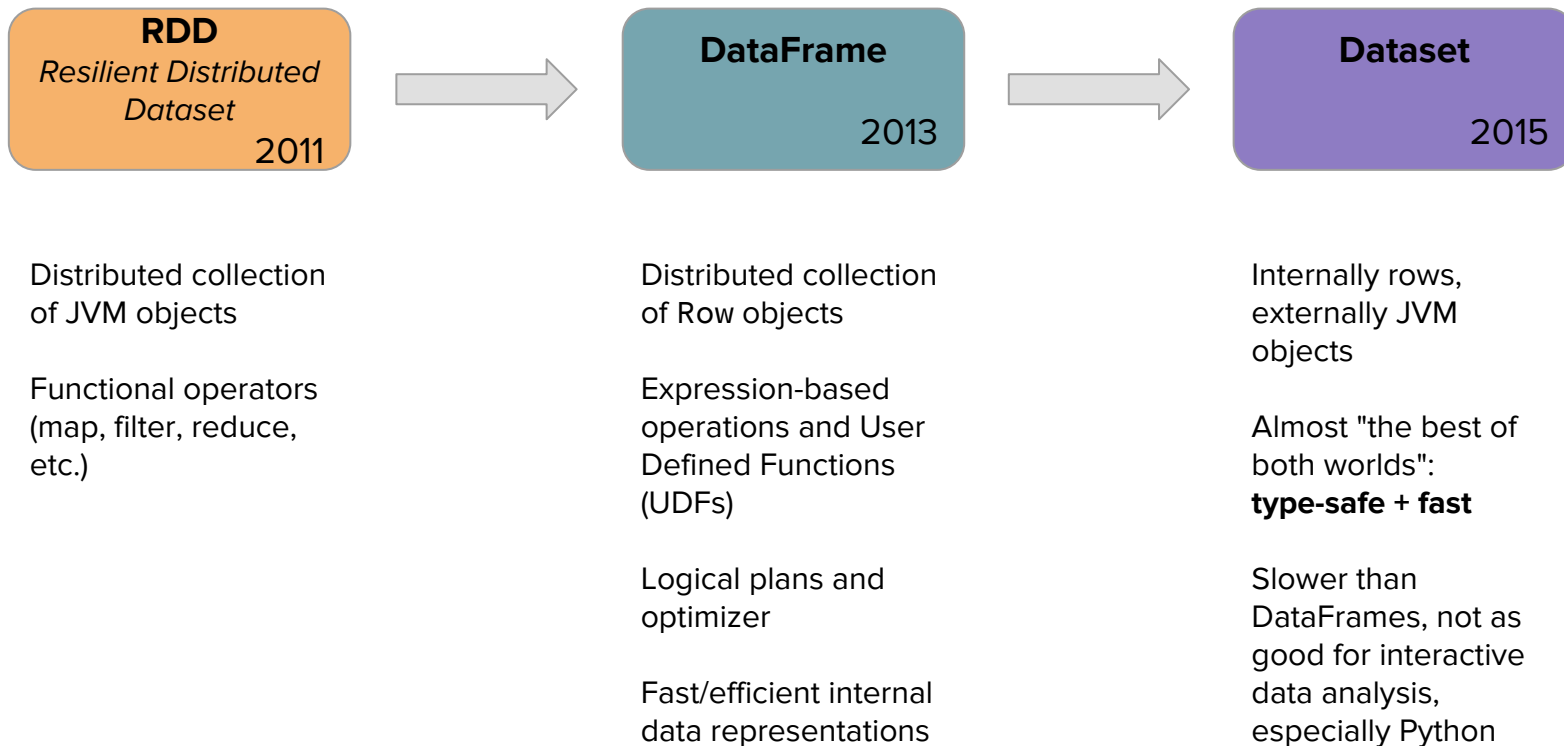


Evolution of Spark APIs



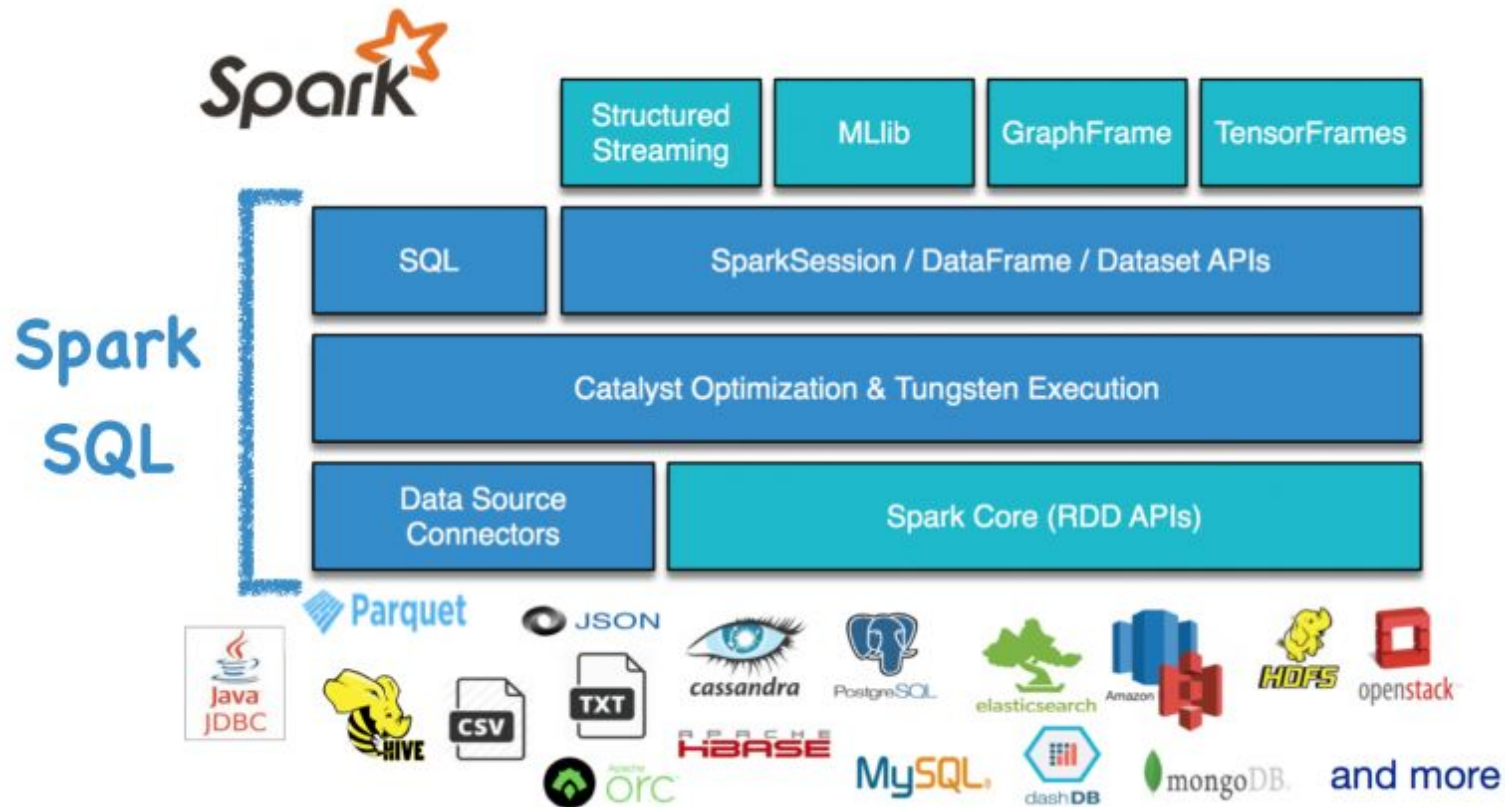
Architecture

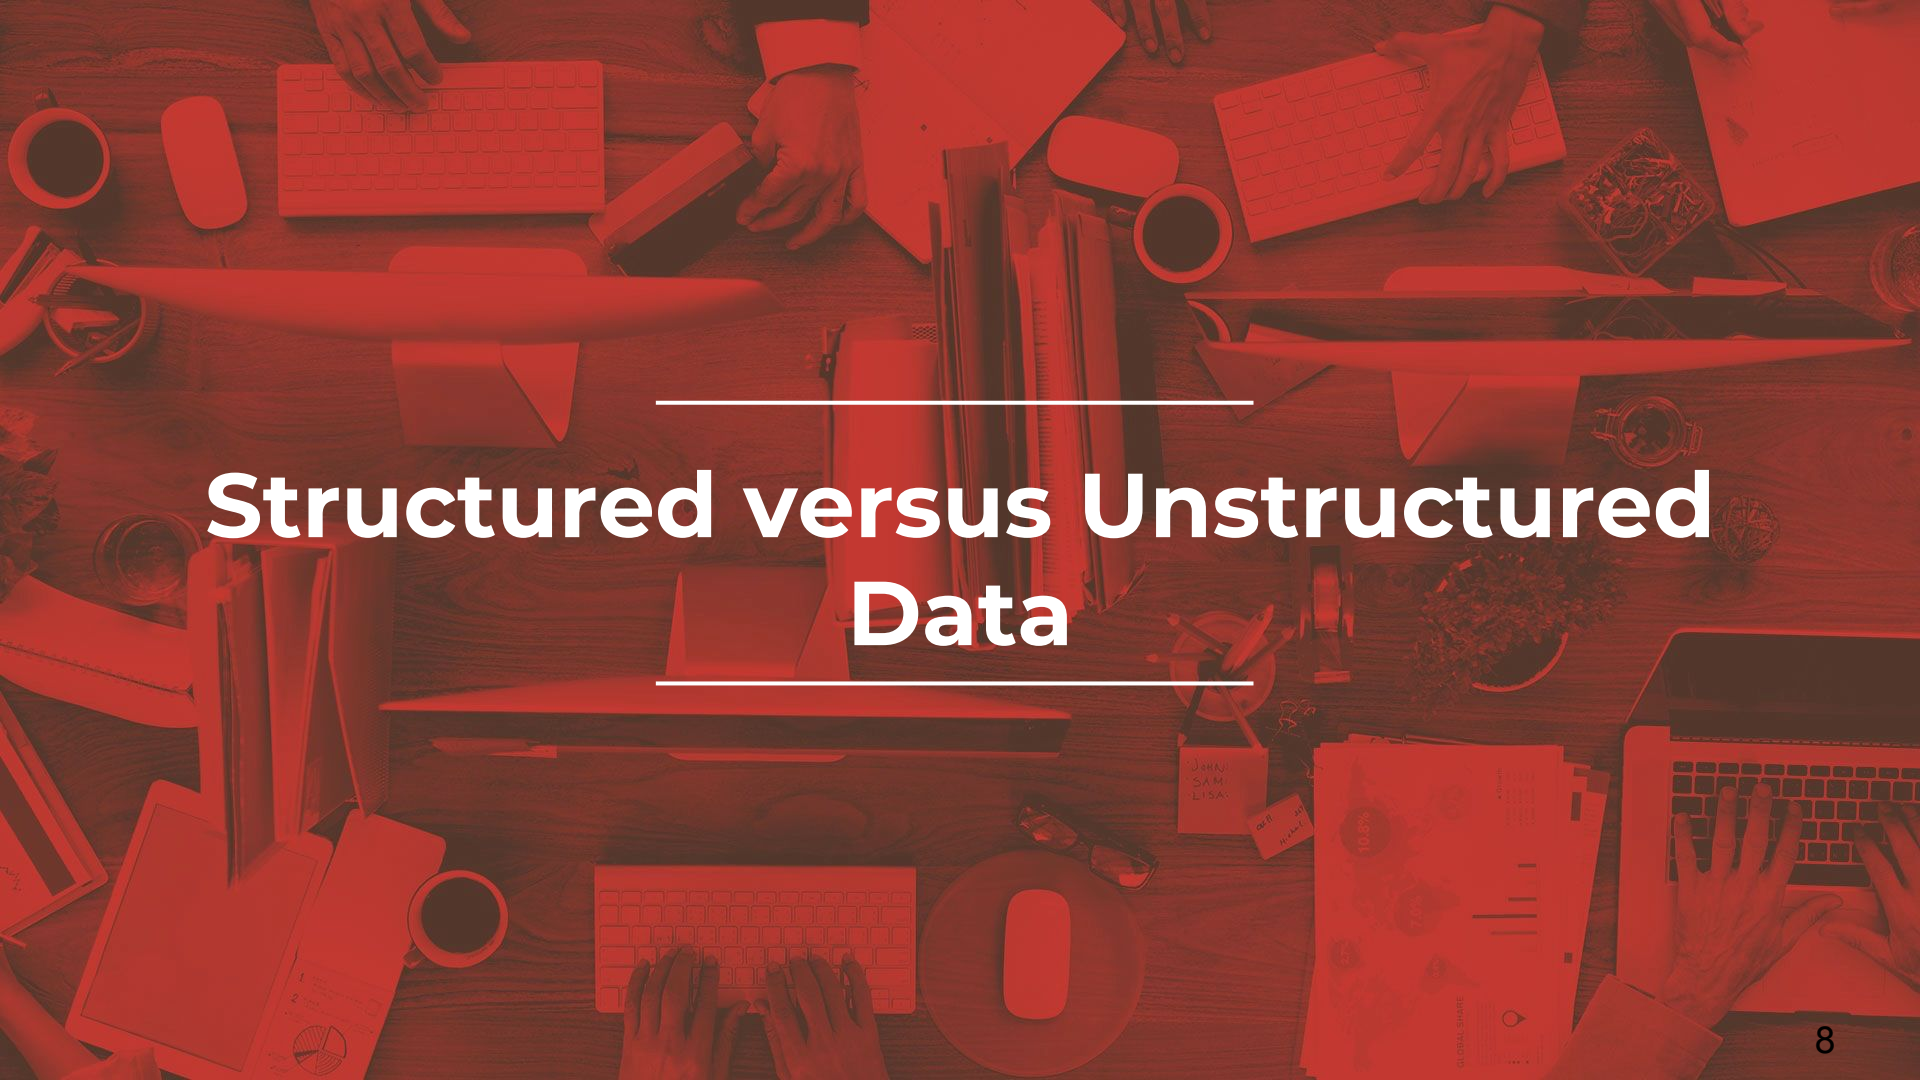
Evolution of Spark APIs





Spark SQL, DataFrames and Datasets

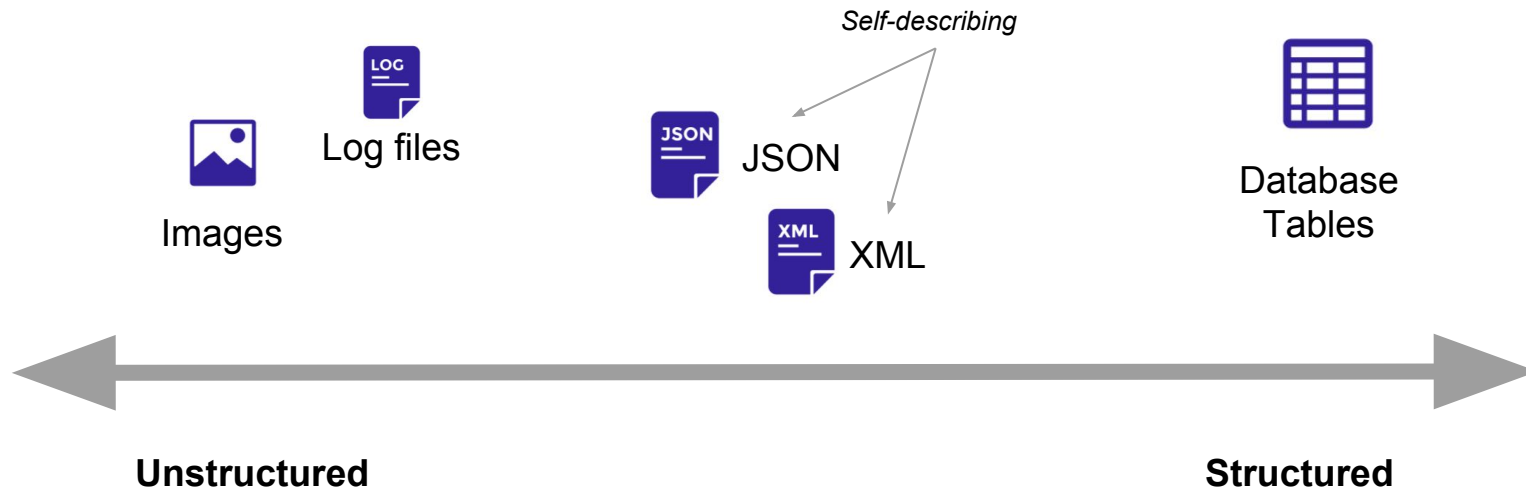




Structured versus Unstructured Data



Structured versus Unstructured Data





Unstructured Data on Spark

RDDs

RDDs are not aware of the *schema* of the data.

```
case class Flight(  
  DEST_COUNTRY_NAME: String, ORIGIN_COUNTRY_NAME: String, count: Integer  
)  
  
                                RDD[Flight]
```

```
+-----+-----+-----+  
|DEST_COUNTRY_NAME|ORIGIN_COUNTRY_NAME|count|  
+-----+-----+-----+  
|      United States|      Romania|      1|  
|      United States|      Ireland|    264|  
|      United States|       India|    691|  
|           Egypt|    United States|    241|  
|Equatorial Guinea|    United States|      1|  
|      United States|      Singapore|    251|  
|      United States|      Grenada|    541|  
|      Costa Rica|    United States|   4771|  
|      Senegal|    United States|    291|  
|      United States|    Marshall Islands|    441|  
+-----+-----+-----+  
only showing top 10 rows
```

Spark knows that RDDs may be parameterized with arbitrary types (objects), but it does not know about the structure of the types.

Pros: More flexibility with data types.

Cons: Optimization is hard without knowledge about the structure.



Unstructured Data Computation on Spark

RDDs

Functional transformations are applied to RDDs.

User-defined **function literals** are used in higher-order functions like **map**, **flatMap** and **filter**.

```
rdd.map((x, y) => x+y)
```

Pros: More flexibility with operations.

Cons: Optimization is hard or impossible to be done automatically.



Structured Data

Data is organized into one or more tables (relations).

Tables contain columns (attributes) and rows (records or tuples).

Tables typically represent a collection of objects of a certain type, such as customers or products.

Similar to relational database tables.

DEST_COUNTRY_NAME	ORIGIN_COUNTRY_NAME	count
United States	Romania	1
United States	Ireland	264
United States	India	69
Egypt	United States	24
Equatorial Guinea	United States	1
United States	Singapore	25
Costa Rica	United States	477



Structured Data Computation

Computations can be more rigid,
think about:

SELECT

WHERE

ORDER BY

GROUP BY

COUNT

DEST_COUNTRY_NAME	ORIGIN_COUNTRY_NAME	count
United States	Romania	1
United States	Ireland	264
United States	India	69
Egypt	United States	24
Equatorial Guinea	United States	1
United States	Singapore	25
Costa Rica	United States	477

Pros: Optimization is easier to be done automatically with structured data.

Cons: Less flexibility with computations and data types.



Optimizations (Dataframes and Datasets)

Catalyst

Query optimizer.

Compiles Spark SQL programs down to RDD operations, which can be reordered to reduce the amount of data that must be read and prune unneeded partitioning.

Tungsten

Off-heap serializer.

Provides highly-specialized data encoders, which are column based and off-heap (avoiding garbage collection overhead).



Structured vs Unstructured Data Recap

What We've Learned So Far

- RDDs are suitable to process unstructured data. This gives flexibility of computation, but there is no opportunity for automatic optimization.
- Dataframes and Datasets are better to process structured data. Automatic optimizations are performed by Spark, but data and computation are more rigid.



Q&A





DataFrames



DataFrames

- A DataFrame is a Dataset organized into named columns.
- Conceptually similar to a table in a relational database.
- Optimizations occur under the hood.
- A DataFrame is represented by a Dataset of Rows.
- Have or Require a known schema.
- Albeit, they are **untyped**.
- Transformations on DataFrames are known as **untyped transformations**.



DataFrames

Creating DataFrames

DataFrames can be created from an existing **RDD** with **.toDF**

```
val flight_data_RDD = spark.sparkContext.parallelize(Seq(  
  ("United States", "Romania", "1"),  
  ("United States", "Ireland", "264"),  
  ("United States", "Singapore", "25")))
```

```
flight_data_RDD: org.apache.spark.rdd.RDD[(String, String, String)] = ParallelCollectionRDD[29] at parallelize
```

```
val flight_data_DF = flight_data_RDD.toDF("DEST_COUNTRY_NAME", "ORIGIN_COUNTRY_NAME", "count")
```

```
flight_data_DF: org.apache.spark.sql.DataFrame = [DEST_COUNTRY_NAME: string, ORIGIN_COUNTRY_NAME: string ... 1 more field]
```



DataFrames

Creating DataFrames

DataFrames can be created using `spark.createDataFrame()`

```
import org.apache.spark.sql.Row
import org.apache.spark.sql.types.{StructField, StringType, IntegerType}
val someData = Seq(
  Row("F3", 3),
  Row("F4", 5),
  Row("F5", 8)
)

val someSchema = List(
  StructField("F", StringType, true),
  StructField("n", IntegerType, true)
)

val someDF = spark.createDataFrame(
  spark.sparkContext.parallelize(someData),
  StructType(someSchema)
)
```



DataFrames

Creating DataFrames

DataFrames can be created by reading from a **data source**: JSON, CSV, Parquet, and such.

```
val flight_data = spark
  .read
  .option("inferSchema", "true")
  .option("header", "true")
  .csv("gs://de-training-input/flight-data/*.csv")
```

```
flight_data: org.apache.spark.sql.DataFrame = [DEST_COUNTRY_NAME: string, ORIGIN_COUNTRY_NAME: string ...]
```



DataFrame Transformations

Look Similar to SQL

Transformations on DataFrames are operations which return a DataFrame and their evaluation is lazy.

```
// This import is needed to use the $-notation/  
import spark.implicits._
```

```
// Select routes with more than 1000000 flights  
flight_data.filter($"count" > 100000).show()
```

```
+-----+-----+-----+  
|DEST_COUNTRY_NAME|ORIGIN_COUNTRY_NAME|count|  
+-----+-----+-----+  
|      United States|      United States|348113|  
|      United States|      United States|370002|  
|      United States|      United States|352742|  
|      United States|      United States|343132|  
|      United States|      United States|347452|  
|      United States|      United States|358354|  
+-----+-----+-----+
```

```
// Print the schema in a tree format  
flight_data.printSchema
```

```
root  
 |-- DEST_COUNTRY_NAME: string (nullable = true)  
 |-- ORIGIN_COUNTRY_NAME: string (nullable = true)  
 |-- count: integer (nullable = true)
```

```
// Select only a column  
flight_data.select("DEST_COUNTRY_NAME").show(3)
```

```
+-----+  
|DEST_COUNTRY_NAME|  
+-----+  
|      United States|  
|      United States|  
|      United States|  
+-----+  
only showing top 3 rows
```

Other common transformations: [join](#), [limit](#), [orderBy](#), [where](#), [as](#), [sort](#), [union](#), [drop](#), etc.



Grouping and Aggregating on Dataframes

One of the Most Common Tasks

groupBy

Function that returns a **RelationalGroupedDataset**

RelationalGroupedDataset has several standard aggregation functions, such as count, sum, max, min and avg.

```
titanic_df.groupBy("Sex").avg("Age").show()
```

```
+-----+-----+
|  Sex  | avg(Age)|
+-----+-----+
|female| 27.915708812260537|
|  male| 30.72664459161148|
+-----+-----+
```




Dataframe Operations

Example Dissected: groupBy, count

```
17  val dataframeWay = flight_data
18    .groupBy("DEST_COUNTRY_NAME")
19    .count()
```



$G_i \leftarrow \text{agg-function} (\text{grouped-data}_i)$
 $G_i \leftarrow \text{count} (\text{grouped-data}_i)$
...



Grouping and Aggregating on Dataframes

One of the Most Common Tasks

groupBy with agg (Untyped Transformation)

You can also call agg and a SparkSQL function, such as: .

```
titanic_df.groupBy("Sex").agg(avg("Age"), avg("Fare")).show()
```

```
+-----+-----+-----+
|  Sex|      avg(Age)|      avg(Fare)|
+-----+-----+-----+
|female|27.915708812260537| 44.47981783439487|
|  male| 30.72664459161148|25.523893414211418|
+-----+-----+-----+
```



Dataframes Recap

What We've Learned So Far

- Dataframes have a schema but are untyped.
- You can perform untyped transformations to Dataframes, including aggregations after a `groupBy`.



Q&A





5 minutes





Exercise Time



Exercise

Grouping and Aggregation on DataFrames

Using the Titanic dataset at <gs://de-training-input/titanic/train.csv> , can you calculate the **ground truth** (number of survivors and deaths) registered on the dataset?

The expected answer is an object that contains the **survivor** and **death** count columns:

Example:

```
ground_truth: org.apache.spark.sql.DataFrame = [survivors: bigint, deaths: bigint]
```

```
+-----+-----+
| survivors | deaths |
+-----+-----+
|    350    |    620 |
+-----+-----+
```

Hint: The reported ground truth on the dataset description is not the answer because we are not using the complete dataset.



Titanic Dataset

Dataset Object Creation

Load the dataset

```
val titanic_df = sqlContext.read.format("csv")  
  .option("header", "true").option("inferSchema", "true")  
  .load("gs://de-training-input/titanic/train.csv")
```

} Scala

```
titanic_df = sqlContext.read.format("csv") \  
  .option("header", "true").option("inferSchema", "true") \  
  .load("gs://de-training-input/titanic/train.csv")
```

} Python

Display the Dataframe schema

```
titanic_df.printSchema()
```



Datasets



Datasets

Creating Datasets

Datasets can be created in different ways:

- From an existing **RDD** or **DataFrame**.

```
import spark.implicits._ // .toDS requires this
val myDS = myRDD.toDS
val myDS = myDF.toDS
```

- From a **data source**:
JSON, CSV, Parquet,
and such.

```
val myDS =  
spark.read.json("people.json").as[Person]
```

...

A Class defines the
structure/type

- From common Scala
types.

```
val myDS = List(  
  "Wizeline", "Academy", "Spark").toDS  
val myDS = spark.createDataset(  
  List("Wizeline", "Academy", "Spark"))
```



Typed Columns

On **Datasets** *typed* operations usually act on `TypedColumn`.

To create a `TypedColumn`, call `.as[. . .]` on an *untyped* `Column`.

```
$"price".as[Double]
```



Dataset Transformations

Typed and Untyped Transformations

Untyped transformations

Transformations for DataFrames are available for Datasets too.

Typed transformations

Typed variants of many DataFrame transformations and additional *higher order RDD-like* transformations like `map`, `flatMap`, etc. are also available.



Common Typed Dataset Transformations

Filtering

`filter(pred: T => Boolean): Dataset[T]`

`map[U](f: T => U): Dataset[U]`

Mapping

`flatMap[U](f: T => TraversableOnce[U]):
Dataset[U]`

Distinct

`distinct(): Dataset[T]`

Grouping

`groupByKey[K](f: T => K):
KeyValueGroupedDataset[K, T]`



Grouped Operations on Datasets

KeyValueGroupedDataset

Calling `groupByKey` on a Dataset return a **KeyValueGroupedDataset**.

KeyValueGroupedDataset has several aggregation operations that return **Dataset** or **KeyValueGroupedDataset**



Common KeyValueGroupedDataset Aggregations

mapValues

`mapValues[W](func: (V) => W): KeyValueGroupedDataset[K, W]`

Returns a new KeyValueGroupedDataset where the given function func has been applied to the data.



Common KeyValueGroupedDataset Aggregations

reduceGroups

```
reduceGroups(f: (V, V) => V): Dataset[(K, V)]
```

Reduce each group using a given binary function.

The function must be commutative and associative to ensure a deterministic result.



Common KeyValueGroupedDataset Aggregations

mapGroups

`mapGroups[U](f: (K, Iterator[V]) => U): Dataset[U]`

Applies the given function to each group of data.

The function must be commutative and associative to ensure a deterministic result.

```
val flights_grouped = flight_data_ds
  .groupBy(_.ORIGIN_COUNTRY_NAME)
  .mapGroups{case(k, iter) => (k, iter.map(x => x.count).toArray)}
  .orderBy("_1")
```

`flights_grouped: org.apache.spark.sql.Dataset[(String, Array[Int])]`



Common KeyValueGroupedDataset Aggregations

flatMapGroups

```
flatMapGroups[U](f: (K, Iterator[V]) => TraversableOnce[U]):  
Dataset[U]
```

Applies the given function to each group of data, flattening the results.



Using the General agg Operation

agg

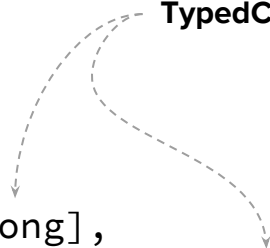
`agg[U](col: TypedColumn[V, U]): Dataset[(K, U)]`

Computes the given aggregation, returning:

- A Dataset of tuples for each unique key.
- The result of computing this aggregation over all elements in the group.

```
val flights_ag = flight_data_ds.groupByKey(
    flight => flight.ORIGIN_COUNTRY_NAME
)
  .agg(
    sum("count").alias("total_flights_from").as[Long],
    max("count").as[Double].alias("max_flights_from").as[Double]
  )
  .orderBy(desc("max_flights_from")).show(10)
```

TypedColumns





Using Aggregator

Aggregator

A base class for user-defined aggregations, which can be used in Dataset operations to take all of the elements of a group and reduce them to a single value.

```
import org.apache.spark.sql.expressions.Aggregator
import org.apache.spark.sql.{Encoder, Encoders}

case class Data(i: Int)

val customSummer = new Aggregator[Data, Int, Int] {
  def zero: Int = 0
  def reduce(b: Int, a: Data): Int = b + a.i
  def merge(b1: Int, b2: Int): Int = b1 + b2
  def finish(r: Int): Int = r
  def bufferEncoder: Encoder[Int] = Encoders.scalaInt
  def outputEncoder: Encoder[Int] = Encoders.scalaInt
}.toColumn

val ds = List(Data(3), Data(5), Data(8)).toDS
val aggregated = ds.select(customSummer)
```

Extracts an `int` from a specific class and adds them up.



Datasets Recap

What We've Learned So Far

- Datasets have a schema and are also typed.
- You can perform all the untyped transformation (as in Dataframes) and also the typed transformations.
- `groupByKey` helps you leverage the typing system and perform grouped aggregations.



Q&A





Exercise Time



Exercise

Grouping and Aggregation on Datasets

Using the Flights dataset, **can you calculate the total registered flights per **ORIGIN** and the number of total destinations for those places?**

The expected answer is a Dataset that contains the **ORIGIN_COUNTRY_NAME**, **total_flights_from**, and **total_flights_to** columns:

```
flights_ag: org.apache.spark.sql.Dataset[org.apache.spark.sql.Row]
```

Hint: If using Datasets, try using data types different from Int for the typed transformations.



Assignment

To Work on Your Own at Home

A photograph of a laptop and a cup of coffee on a wooden desk, with the text overlaid.

**Dataset and Dataframe
Operations Using Alimazon**



Assignment

Dataset and Dataframe Operations

Using the Alamazon Client Purchase Orders Dataset:

1. Best Selling Hours (ID: best_selling_hours)

Problem Description

Alamazon's Marketing Director has issued a request to identify the “spending behaviour by hour” for all kinds of products and report them sorted in nondecreasing order of gross sales.

In particular, they want to know: the average gross sales as well as the minimum and maximum registered sales of the orders placed by hour. Marketing will be using that information to run campaigns during the identified “peak hours”

Input Dataset

You can find the input dataset at:

`gs://de-training-input/alamazon/200000/client-orders/`



Assignment

Dataset and Dataframe Operations

Expected Output Format

- Your final dataset should contain the information described in the problem statement sorted by **average_spending** (in nondecreasing order).
- The hour column goes from 0 to 23.
- Make sure to round the numbers in **average_spending** to two decimal places (you can use the **round** function to accomplish that; see here for details)

| hour | average_spending | min_spending | max_spending |

Output Dataset

You can write your output dataset to your assigned output bucket following this format:

```
gs://de-training-output-<student-name>/assignment-4-best_selling_hours-<attempt>
```

Where <attempt> is an increasing integer (beginning from 0) that distinguishes the multiple attempts you tried (you need this because you don't have permission to erase contents from the output bucket).



Assignment

Dataset and Dataframe Operations

2. Monthly Discount (ID: monthly_discount)

Problem Description

Every month, Alimazon's marketing team assigns a **10%** discount to the **top ten** products that registered the highest number of sales over the **previous six months**, from the time the discount is calculated.

The team wants to operationalize the report below.

Using the Alimazon Client Purchase Orders Dataset, calculate the new product price on a report which should contain: **product_id**, **total_products_sold**, **total_registered_sales**, **unit_product_price**, **new_unit_product_price**.

Input Dataset

You can find the input dataset at:

`gs://de-training-input/alimazon/200000/client-orders/`



Assignment

Dataset and Dataframe Operations

Expected Output Format

- Your final dataset should contain the information described in the problem statement sorted by **total_products_sold** (in nondecreasing order).
- Make sure to round the numbers in **unit_product_price** and **new_unit_product_price** to two decimal places (you can use the **round** function)

```
| product_id | total_products_sold | total_registered_sales | unit_product_price |  
new_unit_product_price |
```

Output Dataset

You can write your output dataset to your assigned output bucket following this format:

```
gs://de-training-output-<student-name>/assignment-4-monthly_discount-<attempt>
```

Where <attempt> is an increasing integer (beginning from 0) that distinguishes the multiple attempts that you tried.



Assignment

Dataset and Dataframe Operations

3. Client Orders Distribution (ID: client_orders_dist)

Problem Description

Also, the sales department at Alimazon is interested in understanding the distribution of **purchases-by-client**, so we were asked to generate that information.

Note: The distribution of purchases-by-user is a table that contains the number of clients (**clients_count**) that have bought the same number of products (**products_count**).

Input Dataset

You can find the input dataset at:

```
gs://de-training-input/alimazon/200000/client-orders/
```



Assignment

Dataset and Dataframe Operations

Expected Output Format

- Your final dataset should contain the information described in the problem statement sorted by **clients_count** (in nondecreasing order).

| products_count | clients_count |

Output Dataset

You can write your output dataset to your assigned output bucket following this format:

```
gs://de-training-output-<student-name>/assignment-4-client_orders_dist-<attempt>
```

Where <attempt> is an increasing integer (beginning from 0) that distinguishes the multiple attempts that you tried.



Where Can I Get this Presentation?

Slack channel PDF



C4 - Data Engineering Academy

Your feedback is very valuable!

<https://goo.gl/forms/LhMpYc1Pr67dFDbw1>